



[Photo by Alina Grubnyak on Unsplash](#)

Introduction to Deep Learning

Why Deep Learning

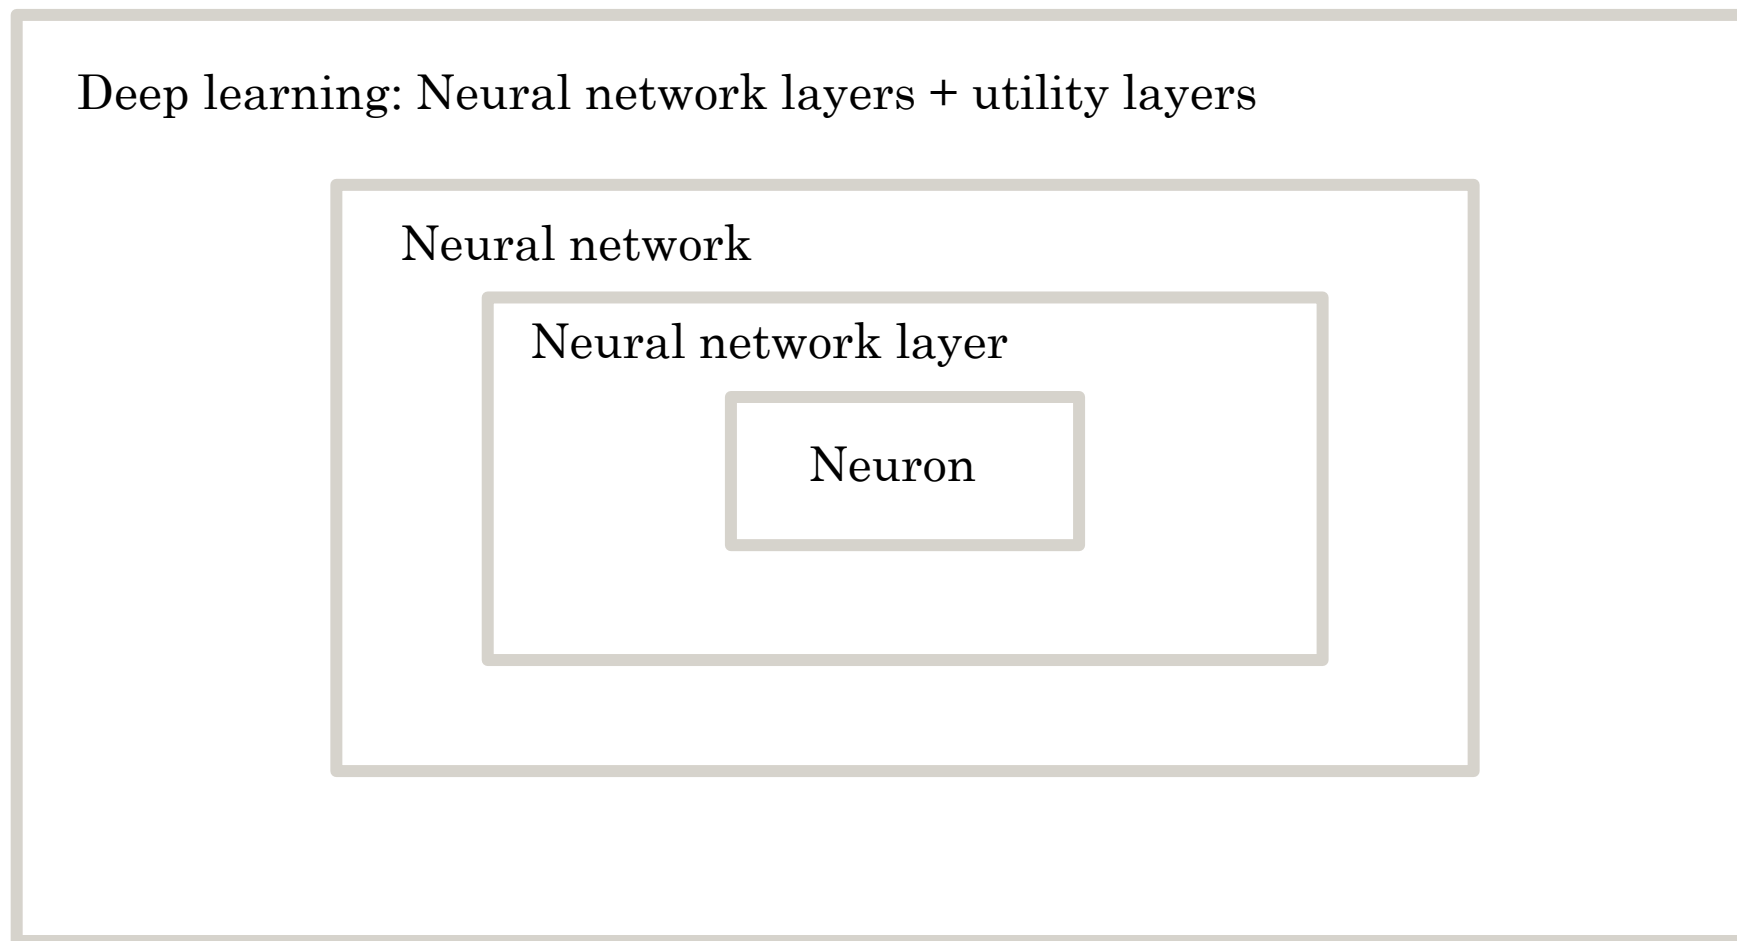
- ✓ A neural network layer is a relatively simple component
- ✓ Easy to scale up by stacking layers
- ✓ Generality: can approximate any function
- ✓ Works well for high dimensions
- ✓ Available hardware optimization
- ✓ Strong community (and open-sourced code)
- X Computationally expensive
- X Can be hard to finetune
- X Hard to interpret (black box)

Yann LeCun

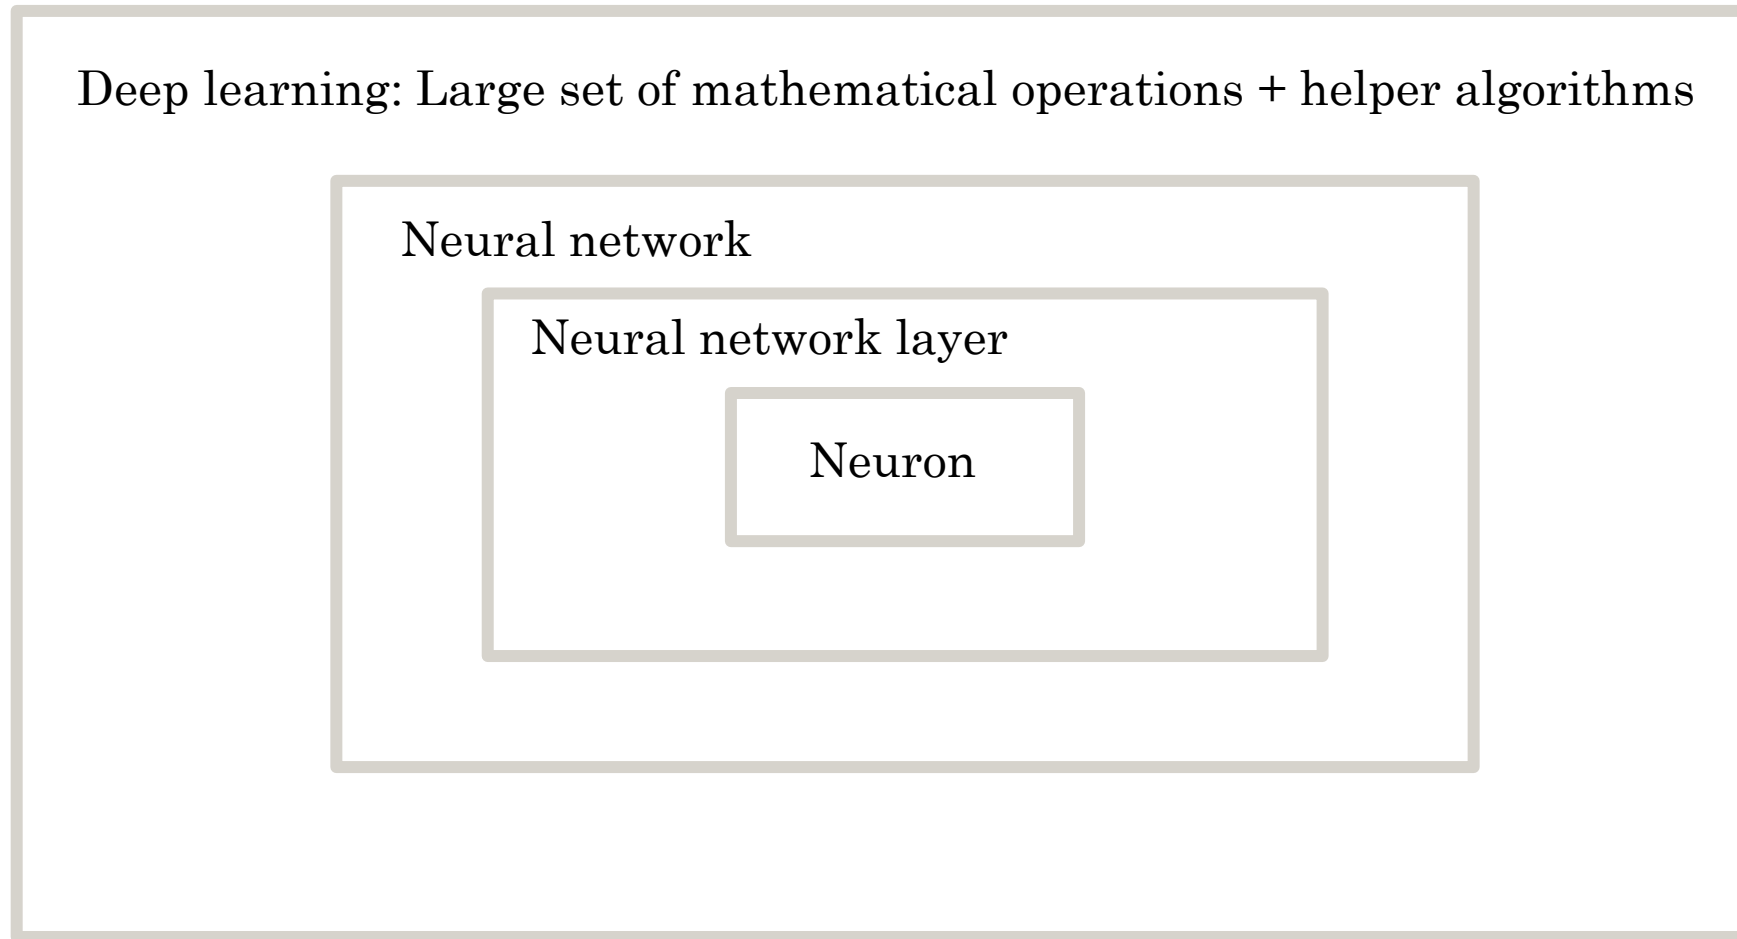


[Wikimedia](#) (retrieved 2023-04-08)

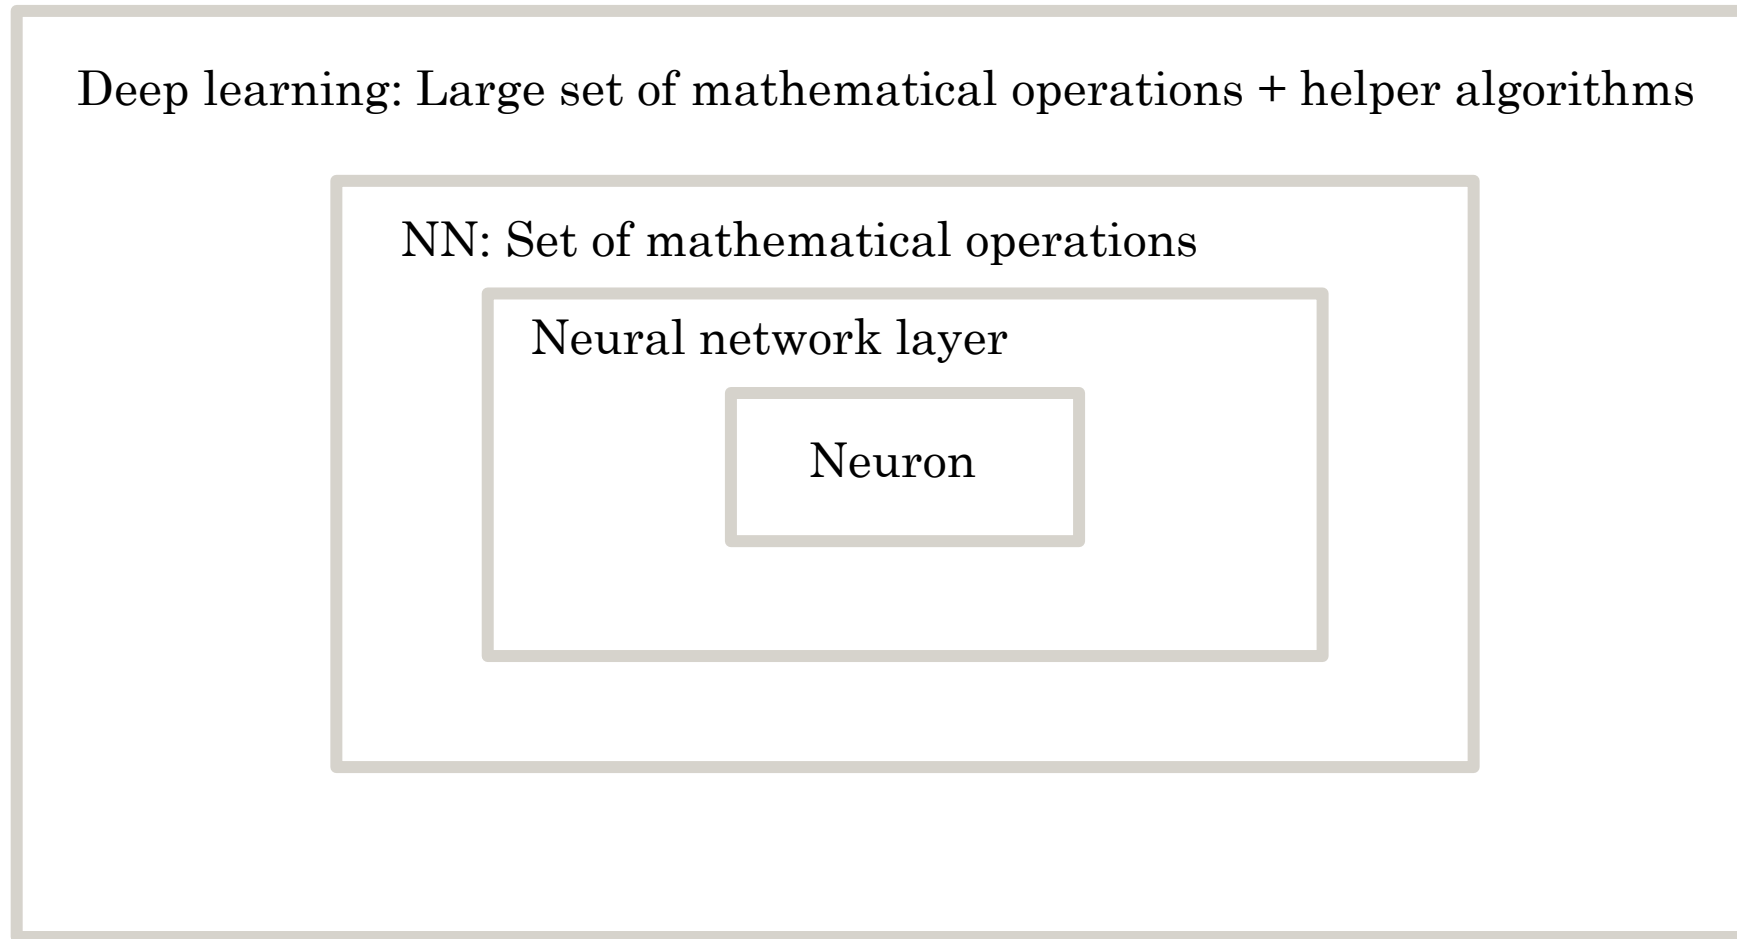
Deep learning components



Deep learning components



Deep learning components



Deep learning components

Deep learning: Large set of mathematical operations + helper algorithms

NN: Set of mathematical operations

NN layer: Set of non-linear functions

Neuron

Deep learning components

Deep learning: Large set of mathematical operations + helper algorithms

NN: Set of mathematical operations

NN layer: Set of non-linear functions

Neuron: non-linear function

Neuron

- A neuron is a function: $z = f(x, w) = f(\sum_i w_i x_i)$

Neuron

- A neuron is a function: $z = f(x, w) = f(\sum_i w_i x_i)$



[Photo by Jan van der Wolf from Pexels](#) (retrieved 2023-04-08)

Neuron

- A neuron is a function: $z = f(x, w) = f(\sum_i w_i x_i)$
- Activation functions:
 - Sigmoid

$$f(x) = \frac{e^x}{e^x + 1}$$



[Photo by Jan van der Wolf from Pexels](#) (retrieved 2023-04-08)

Neuron

- A neuron is a function: $z = f(x, w) = f(\sum_i w_i x_i)$

- Activation functions:

- Sigmoid

$$f(x) = \frac{e^x}{e^x + 1}$$

- ReLU

$$f(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \end{cases}$$



[Photo by Jan van der Wolf from Pexels](#) (retrieved 2023-04-08)

Neuron

- A neuron is a function: $z = f(x, w) = f(\sum_i w_i x_i)$

- Activation functions:

- Sigmoid

$$f(x) = \frac{e^x}{e^x + 1}$$

- ReLU

$$f(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \end{cases}$$

- GELU

$$f(x) = x + P(X \leq x)$$



[Photo by Jan van der Wolf from Pexels](#) (retrieved 2023-04-08)

Neuron

- A neuron is a function: $z = f(x, w) = f(\sum_i w_i x_i)$

- Activation functions:

- Sigmoid

$$f(x) = \frac{e^x}{e^x + 1}$$

- ReLU

$$f(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \end{cases}$$

- GELU


$$f(x) = x + P(X \leq x)$$



CDF of the normal
distribution



[Photo by Jan van der Wolf from Pexels](#) (retrieved 2023-04-08)



[Wikimedia](#) (retrieved 2023-06-25)

Neuron

- In summary, a neuron is a non-linear function:


$$z = f\left(\sum_i w_i x_i\right)$$

$$z_{ReLU} = f\left(\sum_i w_i x_i\right) = \begin{cases} \sum_i w_i x_i & \text{if } \sum_i w_i x_i > 0 \\ 0 & \text{otherwise} \end{cases}$$

$$z_{GELU} = \sum_i w_i x_i + P(X \leq \sum_i w_i x_i)$$



[Photo by Jan van der Wolf from Pexels](#) (retrieved 2023-04-08)



[Wikimedia](#) (retrieved 2023-06-25)

Neuron

- In summary, a neuron is a non-linear function:


$$\mathbf{z} \quad z = f\left(\sum_i w_i x_i\right)$$

$$z_{ReLU} = f\left(\sum_i w_i x_i\right) = \begin{cases} \sum_i w_i x_i & \text{if } \sum_i w_i x_i > 0 \\ 0 & \text{otherwise} \end{cases}$$

$$z_{GELU} = \sum_i w_i x_i + P(X \leq \sum_i w_i x_i)$$



[Photo by Jan van der Wolf from Pexels](#) (retrieved 2023-04-08)




[Wikimedia](#) (retrieved 2023-06-25)

Neural network layer

- In summary, a neuron is a non-linear function:

$$z = f\left(\sum_i^4 w_i x_i\right)$$




[Photo by Jan van der Wolf from Pexels](#) (retrieved 2023-04-08)

Neural network layer

- In summary, a neuron is a non-linear function

$$z = f\left(\sum_i^4 w_i x_i\right)$$

- A NN layer is a set of non-linear functions




[Photo by Jan van der Wolf from Pexels](#) (retrieved 2023-04-08)

Neural network layer

- In summary, a neuron is a non-linear function

$$z_1 = f\left(\sum_i^4 w_i x_i\right)$$

- A NN layer is a set of non-linear functions



[Photo by Jan van der Wolf from Pexels](#) (retrieved 2023-04-08)

Neural network layer


- In summary, a neuron is a non-linear function

$$z_1 = f\left(\sum_i^4 w_i x_i\right)$$

- A NN layer is a set of non-linear functions



[Photo by Jan van der Wolf from Pexels](#) (retrieved 2023-04-08)




Neural network

- A Neural Network is a set of NN layers = set of mathematical operations



[Photo by Jan van der Wolf from Pexels](#) (retrieved 2023-04-08)




Neural network

- A Neural Network is a set of NN layers = set of mathematical operations



[Photo by Jan van der Wolf from Pexels](#) (retrieved 2023-04-08)



Neural network

- A Neural Network is a set of NN layers = set of mathematical operations




[Photo by Jan van der Wolf from Pexels](#) (retrieved 2023-04-08)

Output layer

Hidden layer

Input layer




NN Generality

Learning / training

$$E_B(w) = \sum_{i \in B_k} |y_i - \hat{y}(w)_i|$$

- Goal: find the **weights** w that **minimize** the **error** function.




Learning / training

1. Initialize the weights randomly
 2. Calculate the gradient of the error function
 3. Update weights
- Repeat steps 2 and 3 N times

SGD

1. $w_{j,0} = \text{random}$
2. $v_{j,t} = L \nabla_{w_j} E_B(w_{j,t})$
3. $w_{j,t+1} = w_{j,t} - v_{j,t}$




Learning / training

1. Initialize the weights randomly
 2. Calculate the gradient of the error function
 3. Update weights
- Repeat steps 2 and 3 N times

SGD

1. $w_{j,0} = \text{random}$
2. $v_{j,t} = L \nabla_{w_j} E_B(w_{j,t})$
3. $w_{j,t+1} = \overline{w_{j,t}} - v_{j,t}$



?

Backpropagation

SGD

1. $w_{j,0} = \text{random}$
2. $v_{j,t} = L \nabla_{w_j} E_B(w_{j,t})$
3. $w_{j,t+1} = w_{j,t} - v_{j,t}$


- How do we calculate the derivative of the error function?
- Chain rule

$$E_B(w) = \sum_{i \in B_k} |y_i - \hat{y}(w)_i|$$

$$\Rightarrow E = \frac{1}{2} (y - \hat{y}(w))^2 = \frac{1}{2} \Delta^2$$

$$\begin{aligned} \partial_{w_1} E &= \Delta \partial_{w_1} \hat{y} = \\ \Delta w_9 \partial_{w_1} z_1 &= \\ \Delta w_9 x_1 \end{aligned}$$

$$z = f\left(\sum_i w_i x_i\right)$$



Backpropagation

SGD

1. $w_{j,0} = \text{random}$
2. $v_{j,t} = L \nabla_{w_j} E_B(w_{j,t})$
3. $w_{j,t+1} = w_{j,t} - v_{j,t}$

- How do we calculate the derivative of the error function?
- Chain rule + Matrix multiplication


$$E_B(w) = \sum_{i \in B_k} |y_i - \hat{y}(w)_i|$$

$$\Rightarrow E = \frac{1}{2} (y - \hat{y}(w))^2 = \frac{1}{2} \Delta^2$$

$$z = f\left(\sum_i w_i x_i\right)$$

$$\begin{aligned} \partial_{w_1} E &= \Delta \partial_{w_1} \hat{y} = \\ \Delta w_9 \partial_{w_1} z_1 &= \\ \Delta w_9 x_1 \end{aligned}$$

With many weights and layers
 \Rightarrow use **matrix multiplication**



Backpropagation

SGD

1. $w_{j,0} = \text{random}$
2. $v_{j,t} = L \nabla_{w_j} E_B(w_{j,t})$
3. $w_{j,t+1} = w_{j,t} - v_{j,t}$

- How do we calculate the derivative of the error function?
- Chain rule + Matrix multiplication + storing the intermediate results


$$E_B(w) = \sum_{i \in B_k} |y_i - \hat{y}(w)_i|$$

$$\Rightarrow E = \frac{1}{2} (y - \hat{y}(w))^2 = \frac{1}{2} \Delta^2$$

$$z = f\left(\sum_i w_i x_i\right)$$

$$\begin{aligned} \partial_{w_1} E &= \Delta \partial_{w_1} \hat{y} = \\ \Delta w_9 \partial_{w_1} z_1 &= \\ \Delta w_9 x_1 \end{aligned}$$

With many weights and layers
 \Rightarrow use **matrix multiplication**



Backpropagation

SGD

1. $w_{j,0} = \text{random}$
2. $v_{j,t} = L \nabla_{w_j} E_B(w_{j,t})$
3. $w_{j,t+1} = \underline{w_{j,t} - v_{j,t}}$


- How do we calculate the derivative of the error function?
- Chain rule + Matrix multiplication + storing the intermediate results
- Think of “loss” (i.e. step 2) and “update” (i.e. step 3) as two agents working together to train the network.

$$E_B(w) = \sum_{i \in B_k} |y_i - \hat{y}(w)_i|$$

$$\Rightarrow E = \frac{1}{2} (y - \hat{y}(w))^2 = \frac{1}{2} \Delta^2$$

$$\begin{aligned} \partial_{w_1} E &= \Delta \partial_{w_1} \hat{y} = \\ \Delta w_9 \partial_{w_1} z_1 &= \\ \Delta w_9 x_1 \end{aligned}$$

$$z = f\left(\sum_i w_i x_i\right)$$



Backpropagation

SGD

1. $w_{j,0} = \text{random}$
2. $v_{j,t} = L \nabla_{w_j} E_B(w_{j,t})$
3. $w_{j,t+1} = \underline{w_{j,t}} - v_{j,t}$

- How do we calculate the derivative of the error function?
- Chain rule + Matrix multiplication + storing the intermediate results
- Think of “loss” (i.e. step 2) and “update” (i.e. step 3) as two agents working together to train the network.

$$E_B(w) = \sum_{i \in B_k} |y_i - \hat{y}(w)_i|$$


$$\Rightarrow E = \frac{1}{2} (y - \hat{y}(w))^2 = \frac{1}{2} \Delta^2$$

$$\partial_{w_1} E = \Delta \partial_{w_1} \hat{y} =$$

$$\Delta w_9 \partial_{w_1} z_1 =$$


$$\Delta w_9 x_1$$

$$z = f\left(\sum_i w_i x_i\right)$$



Code example: forward

Implementation in the “real-world”: **PyTorch**



Define the Class


We define our neural network by subclassing `nn.Module`, and initialize the neural network layers in `__init__`. Every `nn.Module` subclass implements the operations on input data in the `forward` method.

```
[ ] class NeuralNetwork(nn.Module):
    def __init__(self):
        super().__init__()
        self.flatten = nn.Flatten()
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(28*28, 512),
            nn.ReLU(),
            nn.Linear(512, 512),
            nn.ReLU(),
            nn.Linear(512, 10),
        )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits
```


Code example: forward

Implementation in the “real-world”: **PyTorch**

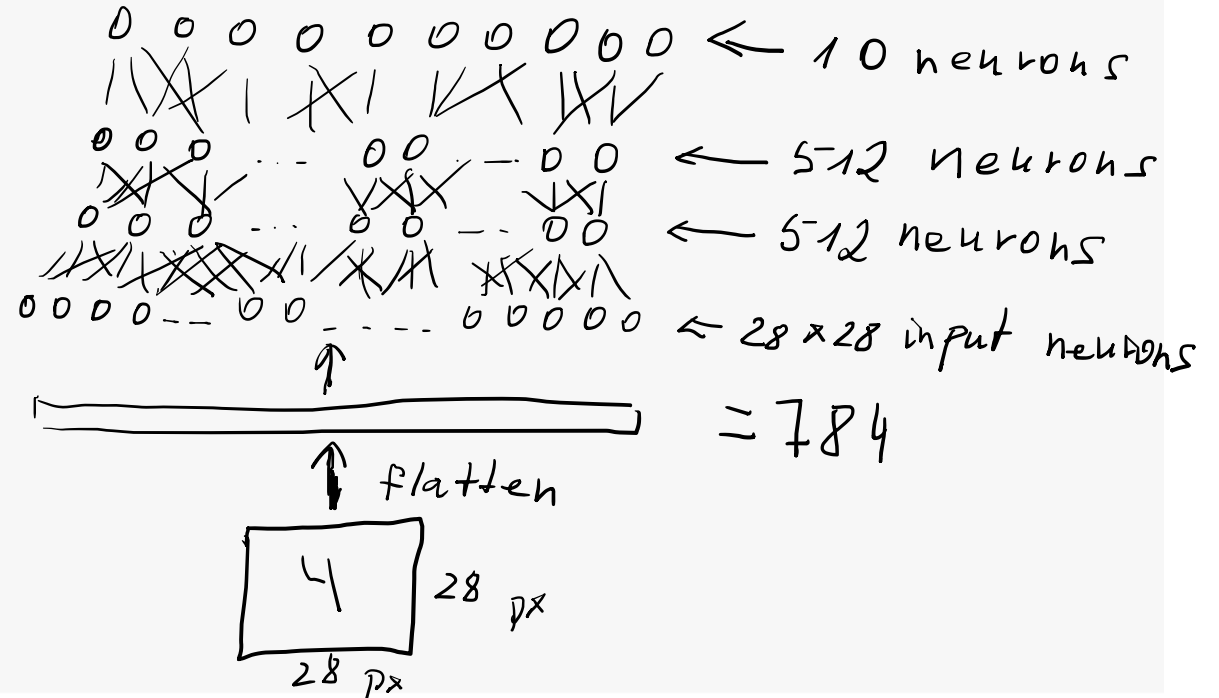


Define the Class

We define our neural network by subclassing `nn.Module`, and initialize the neural network layers in `__init__`. Every `nn.Module` subclass implements the operations on input data in the `forward` method.


```
[ ] class NeuralNetwork(nn.Module):
    def __init__(self):
        super().__init__()
        self.flatten = nn.Flatten()
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(28*28, 512),
            nn.ReLU(),
            nn.Linear(512, 512),
            nn.ReLU(),
            nn.Linear(512, 10),
        )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits
```



Code example: forward

Implementation in the “real-world”: **PyTorch**




To use the model, we pass it the input data. This executes the model's `forward`, along with some [background operations](#). Do not call `model.forward()` directly!

Calling the model on the input returns a 2-dimensional tensor with `dim=0` corresponding to each output of 10 raw predicted values for each class, and `dim=1` corresponding to the individual values of each output. We get the prediction probabilities by passing it through an instance of the `nn.Softmax` module.

```
[ ] X = torch.rand(1, 28, 28, device=device)
    logits = model(X)
    pred_probab = nn.Softmax(dim=1)(logits)
    y_pred = pred_probab.argmax(1)
    print(f"Predicted class: {y_pred}")
```

Code example: backward

Implementation in the “real-world”: PyTorch



Consider the simplest one-layer neural network, with input x , parameters w and b , and some loss function. It can be defined in PyTorch in the following manner:

```
[2] import torch

x = torch.ones(5) # input tensor
y = torch.zeros(3) # expected output
w = torch.randn(5, 3, requires_grad=True)
b = torch.randn(3, requires_grad=True)
z = torch.matmul(x, w)+b
loss = torch.nn.functional.binary_cross_entropy_with_logits(z, y)
```


Computing Gradients

To optimize weights of parameters in the neural network, we need to compute the derivatives of our loss function with respect to parameters, namely, we need $\frac{\partial \text{loss}}{\partial w}$ and $\frac{\partial \text{loss}}{\partial b}$ under some fixed values of x and y . To compute those derivatives, we call `loss.backward()`, and then retrieve the values from `w.grad` and `b.grad`:

```
[3] loss.backward()
    print(w.grad)
    print(b.grad)
```

Code example: backward


Implementation in the “real-world”: PyTorch



Consider the simplest one-layer neural network, with input x , parameters w and b , and some loss function. It can be defined in PyTorch in the following manner:

```
[2] import torch

x = torch.ones(5) # input tensor
y = torch.zeros(3) # expected output
w = torch.randn(5, 3, requires_grad=True)
b = torch.randn(3, requires_grad=True)
z = torch.matmul(x, w)+b
loss = torch.nn.functional.binary_cross_entropy_with_logits(z, y)
```



Computing Gradients

To optimize weights of parameters in the neural network, we need to compute the derivatives of our loss function with respect to parameters, namely, we need $\frac{\partial \text{loss}}{\partial w}$ and $\frac{\partial \text{loss}}{\partial b}$ under some fixed values of x and y . To compute those derivatives, we call `loss.backward()`, and then retrieve the values from `w.grad` and `b.grad`:


```
[3] loss.backward()
    print(w.grad)
    print(b.grad)
```

Code example: full implementation

Implementation in the “real-world”: **PyTorch**

```
[ ] # Initialize the loss function
    loss_fn = nn.CrossEntropyLoss()
    # Initialize the optimizer
    optimizer = torch.optim.SGD(model.parameters(), lr=1e-3)

    epochs = 10
    for t in range(epochs):
        train_loop(train_dataloader, model, loss_fn, optimizer)
        test_loop(test_dataloader, model, loss_fn)
    print("Done!")
```



$$E = - \sum_i y_i \log(f) (1 - y_i) \log(1 - f)$$


Cross-entropy loss

1. $w_0 = \text{random}$
2. $v_t = L \nabla_w E_B(w_t)$
3. $w_{t+1} = w_t - v_t$

SGD optimizer

Code example: full implementation

Implementation in the “real-world”: **PyTorch**



$$E = - \sum_i y_i \log(f) (1 - y_i) \log(1 - f)$$

Cross-entropy loss

1. $w_{j,0} = \text{random}$
2. $v_{j,t} = L \nabla_{w_j} E_B(w_{j,t})$
3. $w_{j,t+1} = \underline{w_{j,t} - v_{j,t}}$

SGD optimizer

```
def train_loop(dataloader, model, loss_fn, optimizer):
    size = len(dataloader.dataset)
    for batch, (X, y) in enumerate(dataloader):
        # Compute prediction and loss
        pred = model(X)
        loss = loss_fn(pred, y)

        # Backpropagation
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

def test_loop(dataloader, model, loss_fn):
    size = len(dataloader.dataset)
    num_batches = len(dataloader)
    test_loss = 0

    with torch.no_grad():
        for X, y in dataloader:
            pred = model(X)
            test_loss += loss_fn(pred, y).item()

    test_loss /= num_batches
    print(f"Avg loss: {test_loss:>8f} \n")
```

Deep learning components

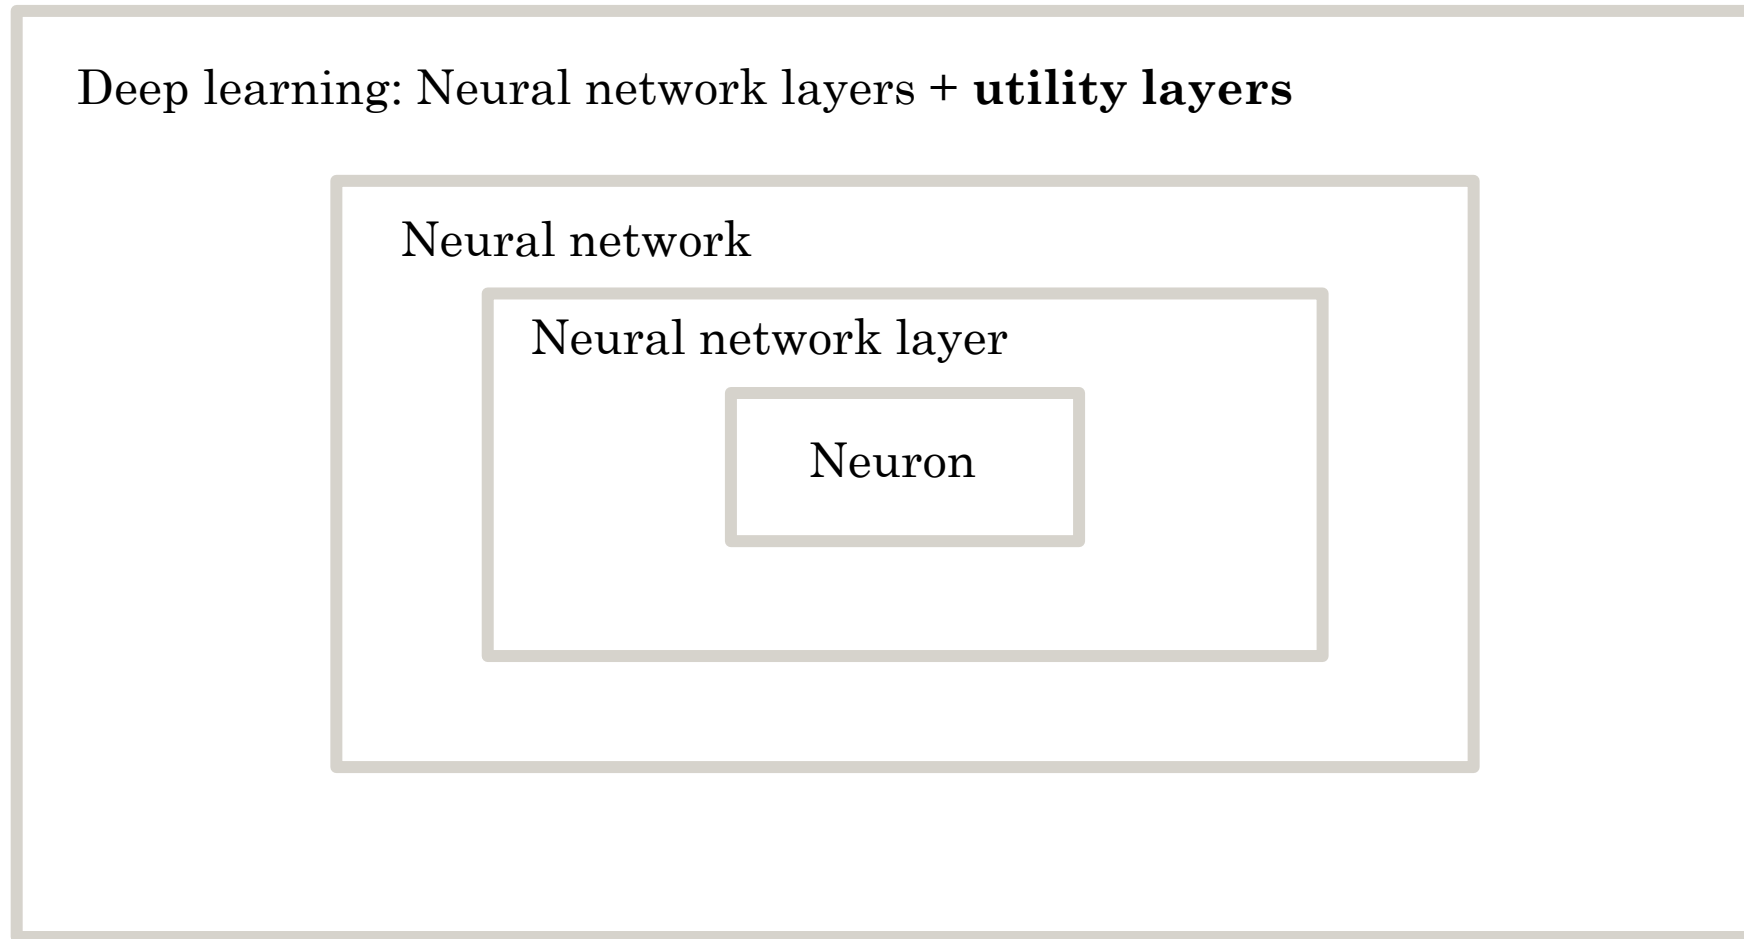
Deep learning: Large set of mathematical operations + helper algorithms

NN: Set of mathematical operations

NN layer: Set of non-linear functions

Neuron: non-linear function


Deep learning components




Dropout

Deep learning: Neural network layers + **utility layers**

Standard
Neural Net




After applying
Dropout



Batch normalization

Deep learning: Neural network layers + **utility layers**

- Training speed-up AND regularization
- Prevents neurons from saturating and gradients from vanishing in deep nets
- Prevents changes in lower layers to significantly impact higher layers



Batch normalization

Deep learning: Neural network layers + **utility layers**

- Training speed-up AND regularization
- Prevents neurons from saturating and gradients from vanishing in deep nets
- Prevents changes in lower layers to significantly impact higher layers

